

```
In [1]: # Convolutional Neural Network

# Installing Theano
# pip install --upgrade --no-deps git+git://github.com/Theano/Theano.git

# Installing Tensorflow
# Install Tensorflow from the website: https://www.tensorflow.org/versions/r0.12/get_started/os_setup.html

# Installing Keras
# pip install --upgrade keras

# Part 1 - Building the CNN

# Importing the Keras Libraries and packages
import numpy as np
import os
import keras_metrics
from keras.models import Sequential
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import TimeDistributed
from keras.layers import LSTM
from keras.layers import Reshape

import warnings
warnings.filterwarnings('ignore')

# Initialising the CNN
classifier = Sequential()

# Step 1 - Convolution
classifier.add(Convolution2D(64, (3, 3), padding = 'same', input_shape = (128, 128, 3), activation = 'relu'))

# Step 2 - Pooling
classifier.add(MaxPooling2D(pool_size = (2, 2)))

# Adding a second convolutional layer
classifier.add(Convolution2D(64, (3, 3), padding = 'same', activation = 'relu'))
classifier.add(MaxPooling2D(pool_size = (2, 2)))

# Adding a third convolutional layer
classifier.add(Convolution2D(64, (3, 3), padding = 'same', activation = 'relu'))
classifier.add(MaxPooling2D(pool_size = (2, 2)))

# Step 3 - Flattening
classifier.add(Flatten())
classifier.add(Dropout(rate = 0.5))

# Step 4 - Full connection
```

```

classifier.add(Dense(output_dim = 128, activation = 'relu'))
classifier.add(Dropout(rate = 0.5))
classifier.add(Dense(output_dim = 7, activation = 'softmax'))

```

```

classifier.summary()

```

Z:\Anaconda3\lib\site-packages\h5py\\_\_init\_\_.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.

from .\_conv import register\_converters as \_register\_converters  
Using TensorFlow backend.

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 128, 128, 64)	1792
max_pooling2d_1 (MaxPooling2)	(None, 64, 64, 64)	0
conv2d_2 (Conv2D)	(None, 64, 64, 64)	36928
max_pooling2d_2 (MaxPooling2)	(None, 32, 32, 64)	0
conv2d_3 (Conv2D)	(None, 32, 32, 64)	36928
max_pooling2d_3 (MaxPooling2)	(None, 16, 16, 64)	0
flatten_1 (Flatten)	(None, 16384)	0
dropout_1 (Dropout)	(None, 16384)	0
dense_1 (Dense)	(None, 128)	2097280
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 7)	903
Total params: 2,173,831		
Trainable params: 2,173,831		
Non-trainable params: 0		

```

In [2]: # Compiling the CNN
classifier.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy', keras_metrics.precision(), keras_metrics.recall()])

```

```
In [3]: # Part 2 - Fitting the CNN to the images

from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale = 1./255,
                                   shear_range = 0.2,
                                   zoom_range = 0.2,
                                   height_shift_range = 0.1,
                                   width_shift_range = 0.1,
                                   channel_shift_range = 10)

test_datagen = ImageDataGenerator(rescale = 1./255)

training_set = train_datagen.flow_from_directory('train/',
                                                target_size = (128, 128),
                                                batch_size = 32,
                                                class_mode = 'categorical')

test_set = test_datagen.flow_from_directory('test/',
                                            target_size = (128, 128),
                                            batch_size = 32,
                                            class_mode = 'categorical')
```

Found 4410 images belonging to 7 classes.  
Found 1475 images belonging to 7 classes.

```
In [4]: results = classifier.fit_generator(training_set,
                                         samples_per_epoch = 4410,
                                         nb_epoch = 100,
                                         validation_data = test_set,
                                         nb_val_samples = 1475)
```

Epoch 1/100  
137/137 [=====] - 993s 7s/step - loss: 1.5628 - acc: 0.3929 - precision: 0.5533 - recall: 0.1415 - val\_loss: 1.2095 - val\_acc: 0.5665 - val\_precision: 0.7853 - val\_recall: 0.2206  
Epoch 2/100  
137/137 [=====] - 1266s 9s/step - loss: 1.2770 - acc: 0.5051 - precision: 0.6919 - recall: 0.2861 - val\_loss: 1.0006 - val\_acc: 0.6136 - val\_precision: 0.7737 - val\_recall: 0.3900  
Epoch 3/100  
137/137 [=====] - 1831s 13s/step - loss: 1.1648 - acc: 0.5520 - precision: 0.7100 - recall: 0.3584 - val\_loss: 0.9996 - val\_acc: 0.6202 - val\_precision: 0.7643 - val\_recall: 0.3972  
Epoch 4/100  
137/137 [=====] - 1842s 13s/step - loss: 1.0854 - acc: 0.5749 - precision: 0.7181 - recall: 0.3954 - val\_loss: 0.7919 - val\_acc: 0.7193 - val\_precision: 0.8122 - val\_recall: 0.5350  
Epoch 5/100  
137/137 [=====] - 1834s 13s/step - loss: 1.0322 - acc: 0.6031 - precision: 0.7268 - recall: 0.4306 - val\_loss: 0.8408 - val\_acc: 0.6862 - val\_precision: 0.7572 - val\_recall: 0.5555  
Epoch 6/100  
137/137 [=====] - 1831s 13s/step - loss: 0.9928 - acc: 0.6052 - precision: 0.7278 - recall: 0.4482 - val\_loss: 0.6830 - val\_acc: 0.7664 - val\_precision: 0.8284 - val\_recall: 0.6355  
Epoch 7/100  
137/137 [=====] - 1829s 13s/step - loss: 0.9311 - acc: 0.6385 - precision: 0.7366 - recall: 0.4946 - val\_loss: 0.6259 - val\_acc: 0.7706 - val\_precision: 0.8500 - val\_recall: 0.6745  
Epoch 8/100  
137/137 [=====] - 1846s 13s/step - loss: 0.8966 - acc: 0.6524 - precision: 0.7515 - recall: 0.5258 - val\_loss: 0.6064 - val\_acc: 0.7807 - val\_precision: 0.8615 - val\_recall: 0.7008  
Epoch 9/100  
137/137 [=====] - 1850s 14s/step - loss: 0.8613 - acc: 0.6698 - precision: 0.7572 - recall: 0.5427 - val\_loss: 0.5174 - val\_acc: 0.8174 - val\_precision: 0.8649 - val\_recall: 0.7658  
Epoch 10/100  
137/137 [=====] - 1849s 13s/step - loss: 0.8137 - acc: 0.6851 - precision: 0.7670 - recall: 0.5757 - val\_loss: 0.5954 - val\_acc: 0.7586 - val\_precision: 0.8229 - val\_recall: 0.6929  
Epoch 11/100  
137/137 [=====] - 1860s 14s/step - loss: 0.8313 - acc: 0.6821 - precision: 0.7609 - recall: 0.5623 - val\_loss: 0.4963 - val\_acc: 0.8127 - val\_precision: 0.8716 - val\_recall: 0.7518  
Epoch 12/100  
137/137 [=====] - 1834s 13s/step - loss: 0.7708 - acc: 0.7062 - precision: 0.7819 - recall: 0.6027 - val\_loss: 0.4024 - val\_acc: 0.8697 - val\_precision: 0.9059 - val\_recall: 0.7956  
Epoch 13/100  
137/137 [=====] - 1879s 14s/step - loss: 0.7324 - acc: 0.7144 - precision: 0.7802 - recall: 0.6262 - val\_loss: 0.3662 - val\_acc: 0.8704 - val\_precision: 0.9020 - val\_recall: 0.8367  
Epoch 14/100  
137/137 [=====] - 1854s 14s/step - loss: 0.7525 - acc: 0.7150 - precision: 0.7818 - recall: 0.6143 - val\_loss: 0.4394 - val\_acc: 0.8488 - val\_precision: 0.8949 - val\_recall: 0.7926  
Epoch 15/100

137/137 [=====] - 1839s 13s/step - loss: 0.7074 - acc: 0.7273 - precision: 0.8069 - recall: 0.6438 - val\_loss: 0.3612 - val\_acc: 0.8658 - val\_precision: 0.8963 - val\_recall: 0.8203  
Epoch 16/100  
137/137 [=====] - 1843s 13s/step - loss: 0.6650 - acc: 0.7406 - precision: 0.7971 - recall: 0.6647 - val\_loss: 0.3041 - val\_acc: 0.8888 - val\_precision: 0.9130 - val\_recall: 0.8684  
Epoch 17/100  
137/137 [=====] - 1839s 13s/step - loss: 0.6458 - acc: 0.7455 - precision: 0.8057 - recall: 0.6788 - val\_loss: 0.2857 - val\_acc: 0.8982 - val\_precision: 0.9134 - val\_recall: 0.8724  
Epoch 18/100  
137/137 [=====] - 1852s 14s/step - loss: 0.6596 - acc: 0.7425 - precision: 0.8055 - recall: 0.6769 - val\_loss: 0.3170 - val\_acc: 0.9024 - val\_precision: 0.9438 - val\_recall: 0.8644  
Epoch 19/100  
137/137 [=====] - 1846s 13s/step - loss: 0.6174 - acc: 0.7584 - precision: 0.8117 - recall: 0.6949 - val\_loss: 0.2950 - val\_acc: 0.8976 - val\_precision: 0.9187 - val\_recall: 0.8718  
Epoch 20/100  
137/137 [=====] - 1845s 13s/step - loss: 0.6245 - acc: 0.7597 - precision: 0.8115 - recall: 0.6943 - val\_loss: 0.2457 - val\_acc: 0.9086 - val\_precision: 0.9244 - val\_recall: 0.8937  
Epoch 21/100  
137/137 [=====] - 1846s 13s/step - loss: 0.6116 - acc: 0.7554 - precision: 0.8093 - recall: 0.7010 - val\_loss: 0.2783 - val\_acc: 0.9191 - val\_precision: 0.9481 - val\_recall: 0.8831  
Epoch 22/100  
137/137 [=====] - 1840s 13s/step - loss: 0.5690 - acc: 0.7748 - precision: 0.8231 - recall: 0.7254 - val\_loss: 0.2455 - val\_acc: 0.9235 - val\_precision: 0.9395 - val\_recall: 0.8922  
Epoch 23/100  
137/137 [=====] - 1842s 13s/step - loss: 0.5622 - acc: 0.7764 - precision: 0.8283 - recall: 0.7237 - val\_loss: 0.2486 - val\_acc: 0.9213 - val\_precision: 0.9404 - val\_recall: 0.8976  
Epoch 24/100  
137/137 [=====] - 1829s 13s/step - loss: 0.5844 - acc: 0.7788 - precision: 0.8269 - recall: 0.7225 - val\_loss: 0.2040 - val\_acc: 0.9377 - val\_precision: 0.9577 - val\_recall: 0.9214  
Epoch 25/100  
137/137 [=====] - 1837s 13s/step - loss: 0.5524 - acc: 0.7851 - precision: 0.8332 - recall: 0.7386 - val\_loss: 0.2087 - val\_acc: 0.9390 - val\_precision: 0.9489 - val\_recall: 0.9172  
Epoch 26/100  
137/137 [=====] - 1847s 13s/step - loss: 0.5296 - acc: 0.8049 - precision: 0.8484 - recall: 0.7560 - val\_loss: 0.1894 - val\_acc: 0.9485 - val\_precision: 0.9563 - val\_recall: 0.9322  
Epoch 27/100  
137/137 [=====] - 1849s 13s/step - loss: 0.5140 - acc: 0.8024 - precision: 0.8470 - recall: 0.7590 - val\_loss: 0.2195 - val\_acc: 0.9321 - val\_precision: 0.9442 - val\_recall: 0.9085  
Epoch 28/100  
137/137 [=====] - 1862s 14s/step - loss: 0.5159 - acc: 0.7996 - precision: 0.8408 - recall: 0.7545 - val\_loss: 0.1993 - val\_acc: 0.9220 - val\_precision: 0.9365 - val\_recall: 0.9098  
Epoch 29/100  
137/137 [=====] - 1829s 13s/step - loss: 0.4758 - acc:

c: 0.8172 - precision: 0.8520 - recall: 0.7802 - val\_loss: 0.1845 - val\_acc: 0.9498 - val\_precision: 0.9623 - val\_recall: 0.9348  
Epoch 30/100  
137/137 [=====] - 1840s 13s/step - loss: 0.4909 - acc: 0.8159 - precision: 0.8519 - recall: 0.7757 - val\_loss: 0.2116 - val\_acc: 0.9281 - val\_precision: 0.9430 - val\_recall: 0.9091  
Epoch 31/100  
137/137 [=====] - 1833s 13s/step - loss: 0.4823 - acc: 0.8077 - precision: 0.8475 - recall: 0.7664 - val\_loss: 0.1374 - val\_acc: 0.9592 - val\_precision: 0.9633 - val\_recall: 0.9457  
Epoch 32/100  
137/137 [=====] - 1833s 13s/step - loss: 0.4921 - acc: 0.8140 - precision: 0.8511 - recall: 0.7714 - val\_loss: 0.1876 - val\_acc: 0.9452 - val\_precision: 0.9590 - val\_recall: 0.9343  
Epoch 33/100  
137/137 [=====] - 1837s 13s/step - loss: 0.4763 - acc: 0.8139 - precision: 0.8492 - recall: 0.7771 - val\_loss: 0.1725 - val\_acc: 0.9404 - val\_precision: 0.9488 - val\_recall: 0.9295  
Epoch 34/100  
137/137 [=====] - 1844s 13s/step - loss: 0.4543 - acc: 0.8300 - precision: 0.8616 - recall: 0.7904 - val\_loss: 0.1274 - val\_acc: 0.9591 - val\_precision: 0.9641 - val\_recall: 0.9517  
Epoch 35/100  
137/137 [=====] - 1844s 13s/step - loss: 0.4583 - acc: 0.8238 - precision: 0.8555 - recall: 0.7873 - val\_loss: 0.1422 - val\_acc: 0.9559 - val\_precision: 0.9622 - val\_recall: 0.9484  
Epoch 36/100  
137/137 [=====] - 1840s 13s/step - loss: 0.4286 - acc: 0.8388 - precision: 0.8688 - recall: 0.8096 - val\_loss: 0.1453 - val\_acc: 0.9579 - val\_precision: 0.9641 - val\_recall: 0.9505  
Epoch 37/100  
137/137 [=====] - 1835s 13s/step - loss: 0.4610 - acc: 0.8296 - precision: 0.8604 - recall: 0.7947 - val\_loss: 0.1219 - val\_acc: 0.9641 - val\_precision: 0.9689 - val\_recall: 0.9525  
Epoch 38/100  
137/137 [=====] - 1832s 13s/step - loss: 0.4229 - acc: 0.8409 - precision: 0.8710 - recall: 0.8081 - val\_loss: 0.1069 - val\_acc: 0.9688 - val\_precision: 0.9780 - val\_recall: 0.9613  
Epoch 39/100  
137/137 [=====] - 1843s 13s/step - loss: 0.4381 - acc: 0.8296 - precision: 0.8623 - recall: 0.8006 - val\_loss: 0.1082 - val\_acc: 0.9709 - val\_precision: 0.9741 - val\_recall: 0.9661  
Epoch 40/100  
137/137 [=====] - 1838s 13s/step - loss: 0.4066 - acc: 0.8456 - precision: 0.8738 - recall: 0.8195 - val\_loss: 0.0965 - val\_acc: 0.9750 - val\_precision: 0.9802 - val\_recall: 0.9688  
Epoch 41/100  
137/137 [=====] - 1838s 13s/step - loss: 0.3972 - acc: 0.8460 - precision: 0.8737 - recall: 0.8205 - val\_loss: 0.1012 - val\_acc: 0.9668 - val\_precision: 0.9713 - val\_recall: 0.9628  
Epoch 42/100  
137/137 [=====] - 1841s 13s/step - loss: 0.4016 - acc: 0.8489 - precision: 0.8708 - recall: 0.8204 - val\_loss: 0.0959 - val\_acc: 0.9763 - val\_precision: 0.9803 - val\_recall: 0.9736  
Epoch 43/100  
137/137 [=====] - 1840s 13s/step - loss: 0.4187 - acc: 0.8436 - precision: 0.8707 - recall: 0.8104 - val\_loss: 0.0904 - val\_acc:

0.9783 - val\_precision: 0.9815 - val\_recall: 0.9729  
Epoch 44/100  
137/137 [=====] - 1929s 14s/step - loss: 0.3989 - acc: 0.8504 - precision: 0.8804 - recall: 0.8205 - val\_loss: 0.1351 - val\_acc: 0.9525 - val\_precision: 0.9574 - val\_recall: 0.9457  
Epoch 45/100  
137/137 [=====] - 1849s 13s/step - loss: 0.3986 - acc: 0.8524 - precision: 0.8760 - recall: 0.8245 - val\_loss: 0.0982 - val\_acc: 0.9702 - val\_precision: 0.9721 - val\_recall: 0.9668  
Epoch 46/100  
137/137 [=====] - 1845s 13s/step - loss: 0.3921 - acc: 0.8526 - precision: 0.8809 - recall: 0.8255 - val\_loss: 0.0914 - val\_acc: 0.9688 - val\_precision: 0.9726 - val\_recall: 0.9627  
Epoch 47/100  
137/137 [=====] - 1850s 14s/step - loss: 0.3903 - acc: 0.8495 - precision: 0.8778 - recall: 0.8253 - val\_loss: 0.0883 - val\_acc: 0.9721 - val\_precision: 0.9774 - val\_recall: 0.9694  
Epoch 48/100  
137/137 [=====] - 1853s 14s/step - loss: 0.3890 - acc: 0.8565 - precision: 0.8800 - recall: 0.8348 - val\_loss: 0.0785 - val\_acc: 0.9783 - val\_precision: 0.9822 - val\_recall: 0.9708  
Epoch 49/100  
137/137 [=====] - 1849s 13s/step - loss: 0.3845 - acc: 0.8560 - precision: 0.8803 - recall: 0.8310 - val\_loss: 0.0875 - val\_acc: 0.9735 - val\_precision: 0.9801 - val\_recall: 0.9702  
Epoch 50/100  
137/137 [=====] - 1853s 14s/step - loss: 0.3961 - acc: 0.8518 - precision: 0.8753 - recall: 0.8235 - val\_loss: 0.1179 - val\_acc: 0.9520 - val\_precision: 0.9576 - val\_recall: 0.9466  
Epoch 51/100  
137/137 [=====] - 1860s 14s/step - loss: 0.3862 - acc: 0.8579 - precision: 0.8841 - recall: 0.8314 - val\_loss: 0.0745 - val\_acc: 0.9823 - val\_precision: 0.9843 - val\_recall: 0.9816  
Epoch 52/100  
137/137 [=====] - 1852s 14s/step - loss: 0.3670 - acc: 0.8600 - precision: 0.8814 - recall: 0.8375 - val\_loss: 0.0731 - val\_acc: 0.9791 - val\_precision: 0.9810 - val\_recall: 0.9729  
Epoch 53/100  
137/137 [=====] - 1850s 14s/step - loss: 0.3722 - acc: 0.8620 - precision: 0.8824 - recall: 0.8333 - val\_loss: 0.0807 - val\_acc: 0.9782 - val\_precision: 0.9815 - val\_recall: 0.9735  
Epoch 54/100  
137/137 [=====] - 1849s 13s/step - loss: 0.3593 - acc: 0.8670 - precision: 0.8898 - recall: 0.8440 - val\_loss: 0.0657 - val\_acc: 0.9817 - val\_precision: 0.9843 - val\_recall: 0.9790  
Epoch 55/100  
137/137 [=====] - 1854s 14s/step - loss: 0.3166 - acc: 0.8824 - precision: 0.9013 - recall: 0.8609 - val\_loss: 0.0626 - val\_acc: 0.9825 - val\_precision: 0.9838 - val\_recall: 0.9798  
Epoch 56/100  
137/137 [=====] - 1850s 14s/step - loss: 0.3454 - acc: 0.8751 - precision: 0.8934 - recall: 0.8550 - val\_loss: 0.0783 - val\_acc: 0.9784 - val\_precision: 0.9830 - val\_recall: 0.9757  
Epoch 57/100  
137/137 [=====] - 1837s 13s/step - loss: 0.3511 - acc: 0.8672 - precision: 0.8874 - recall: 0.8437 - val\_loss: 0.0744 - val\_acc: 0.9776 - val\_precision: 0.9809 - val\_recall: 0.9756



Epoch 58/100  
137/137 [=====] - 1835s 13s/step - loss: 0.3382 - acc: 0.8714 - precision: 0.8921 - recall: 0.8540 - val\_loss: 0.0694 - val\_acc: 0.9796 - val\_precision: 0.9816 - val\_recall: 0.9776  
Epoch 59/100  
137/137 [=====] - 1847s 13s/step - loss: 0.3088 - acc: 0.8890 - precision: 0.9029 - recall: 0.8685 - val\_loss: 0.0648 - val\_acc: 0.9823 - val\_precision: 0.9830 - val\_recall: 0.9810  
Epoch 60/100  
137/137 [=====] - 1867s 14s/step - loss: 0.3464 - acc: 0.8770 - precision: 0.8957 - recall: 0.8572 - val\_loss: 0.0818 - val\_acc: 0.9750 - val\_precision: 0.9808 - val\_recall: 0.9695  
Epoch 61/100  
137/137 [=====] - 1864s 14s/step - loss: 0.3166 - acc: 0.8800 - precision: 0.8959 - recall: 0.8631 - val\_loss: 0.0647 - val\_acc: 0.9850 - val\_precision: 0.9903 - val\_recall: 0.9803  
Epoch 62/100  
137/137 [=====] - 1849s 13s/step - loss: 0.3161 - acc: 0.8771 - precision: 0.8941 - recall: 0.8631 - val\_loss: 0.0732 - val\_acc: 0.9804 - val\_precision: 0.9843 - val\_recall: 0.9763  
Epoch 63/100  
137/137 [=====] - 1849s 13s/step - loss: 0.3326 - acc: 0.8748 - precision: 0.8954 - recall: 0.8553 - val\_loss: 0.0640 - val\_acc: 0.9898 - val\_precision: 0.9905 - val\_recall: 0.9851  
Epoch 64/100  
137/137 [=====] - 1846s 13s/step - loss: 0.3421 - acc: 0.8722 - precision: 0.8923 - recall: 0.8528 - val\_loss: 0.0645 - val\_acc: 0.9891 - val\_precision: 0.9891 - val\_recall: 0.9844  
Epoch 65/100  
137/137 [=====] - 1835s 13s/step - loss: 0.3161 - acc: 0.8804 - precision: 0.8979 - recall: 0.8568 - val\_loss: 0.0696 - val\_acc: 0.9803 - val\_precision: 0.9836 - val\_recall: 0.9776  
Epoch 66/100  
137/137 [=====] - 1845s 13s/step - loss: 0.2990 - acc: 0.8865 - precision: 0.9012 - recall: 0.8719 - val\_loss: 0.0623 - val\_acc: 0.9844 - val\_precision: 0.9864 - val\_recall: 0.9823  
Epoch 67/100  
137/137 [=====] - 1853s 14s/step - loss: 0.3316 - acc: 0.8757 - precision: 0.8982 - recall: 0.8515 - val\_loss: 0.0534 - val\_acc: 0.9878 - val\_precision: 0.9891 - val\_recall: 0.9844  
Epoch 68/100  
137/137 [=====] - 1849s 13s/step - loss: 0.3215 - acc: 0.8843 - precision: 0.9017 - recall: 0.8663 - val\_loss: 0.0644 - val\_acc: 0.9838 - val\_precision: 0.9858 - val\_recall: 0.9817  
Epoch 69/100  
137/137 [=====] - 1854s 14s/step - loss: 0.3186 - acc: 0.8841 - precision: 0.9019 - recall: 0.8686 - val\_loss: 0.0690 - val\_acc: 0.9803 - val\_precision: 0.9809 - val\_recall: 0.9770  
Epoch 70/100  
137/137 [=====] - 1850s 14s/step - loss: 0.3057 - acc: 0.8885 - precision: 0.9044 - recall: 0.8714 - val\_loss: 0.0539 - val\_acc: 0.9837 - val\_precision: 0.9850 - val\_recall: 0.9810  
Epoch 71/100  
137/137 [=====] - 1855s 14s/step - loss: 0.2995 - acc: 0.8885 - precision: 0.9065 - recall: 0.8722 - val\_loss: 0.0572 - val\_acc: 0.9824 - val\_precision: 0.9830 - val\_recall: 0.9790  
Epoch 72/100

137/137 [=====] - 1856s 14s/step - loss: 0.3089 - acc: 0.8816 - precision: 0.8969 - recall: 0.8653 - val\_loss: 0.0648 - val\_acc: 0.9830 - val\_precision: 0.9850 - val\_recall: 0.9823  
Epoch 73/100  
137/137 [=====] - 1857s 14s/step - loss: 0.2994 - acc: 0.8876 - precision: 0.9056 - recall: 0.8686 - val\_loss: 0.0364 - val\_acc: 0.9932 - val\_precision: 0.9939 - val\_recall: 0.9919  
Epoch 74/100  
137/137 [=====] - 1842s 13s/step - loss: 0.2925 - acc: 0.8971 - precision: 0.9078 - recall: 0.8796 - val\_loss: 0.0574 - val\_acc: 0.9864 - val\_precision: 0.9891 - val\_recall: 0.9844  
Epoch 75/100  
137/137 [=====] - 1851s 14s/step - loss: 0.3014 - acc: 0.8926 - precision: 0.9105 - recall: 0.8749 - val\_loss: 0.0376 - val\_acc: 0.9926 - val\_precision: 0.9959 - val\_recall: 0.9913  
Epoch 76/100  
137/137 [=====] - 1868s 14s/step - loss: 0.2748 - acc: 0.8982 - precision: 0.9141 - recall: 0.8847 - val\_loss: 0.0484 - val\_acc: 0.9865 - val\_precision: 0.9878 - val\_recall: 0.9844  
Epoch 77/100  
137/137 [=====] - 1852s 14s/step - loss: 0.2952 - acc: 0.8992 - precision: 0.9151 - recall: 0.8840 - val\_loss: 0.0428 - val\_acc: 0.9946 - val\_precision: 0.9952 - val\_recall: 0.9919  
Epoch 78/100  
137/137 [=====] - 1860s 14s/step - loss: 0.2999 - acc: 0.8859 - precision: 0.9030 - recall: 0.8711 - val\_loss: 0.0413 - val\_acc: 0.9946 - val\_precision: 0.9946 - val\_recall: 0.9933  
Epoch 79/100  
137/137 [=====] - 1861s 14s/step - loss: 0.2707 - acc: 0.8941 - precision: 0.9092 - recall: 0.8788 - val\_loss: 0.0449 - val\_acc: 0.9912 - val\_precision: 0.9932 - val\_recall: 0.9898  
Epoch 80/100  
137/137 [=====] - 1854s 14s/step - loss: 0.2842 - acc: 0.8964 - precision: 0.9121 - recall: 0.8807 - val\_loss: 0.0302 - val\_acc: 0.9939 - val\_precision: 0.9953 - val\_recall: 0.9919  
Epoch 81/100  
137/137 [=====] - 1868s 14s/step - loss: 0.2727 - acc: 0.9025 - precision: 0.9179 - recall: 0.8874 - val\_loss: 0.0315 - val\_acc: 0.9926 - val\_precision: 0.9932 - val\_recall: 0.9905  
Epoch 82/100  
137/137 [=====] - 1882s 14s/step - loss: 0.2918 - acc: 0.8928 - precision: 0.9059 - recall: 0.8755 - val\_loss: 0.0279 - val\_acc: 0.9946 - val\_precision: 0.9966 - val\_recall: 0.9946  
Epoch 83/100  
137/137 [=====] - 1886s 14s/step - loss: 0.2660 - acc: 0.9044 - precision: 0.9139 - recall: 0.8911 - val\_loss: 0.0407 - val\_acc: 0.9878 - val\_precision: 0.9898 - val\_recall: 0.9865  
Epoch 84/100  
137/137 [=====] - 1858s 14s/step - loss: 0.2755 - acc: 0.8998 - precision: 0.9145 - recall: 0.8847 - val\_loss: 0.0325 - val\_acc: 0.9905 - val\_precision: 0.9925 - val\_recall: 0.9878  
Epoch 85/100  
137/137 [=====] - 1852s 14s/step - loss: 0.2845 - acc: 0.8984 - precision: 0.9093 - recall: 0.8832 - val\_loss: 0.0494 - val\_acc: 0.9844 - val\_precision: 0.9884 - val\_recall: 0.9837  
Epoch 86/100  
137/137 [=====] - 1862s 14s/step - loss: 0.2728 - acc:

c: 0.9011 - precision: 0.9149 - recall: 0.8867 - val\_loss: 0.0458 - val\_acc: 0.9844 - val\_precision: 0.9890 - val\_recall: 0.9824  
Epoch 87/100  
137/137 [=====] - 1854s 14s/step - loss: 0.2667 - acc: 0.9036 - precision: 0.9178 - recall: 0.8906 - val\_loss: 0.0341 - val\_acc: 0.9960 - val\_precision: 0.9966 - val\_recall: 0.9960  
Epoch 88/100  
137/137 [=====] - 1852s 14s/step - loss: 0.2757 - acc: 0.8990 - precision: 0.9127 - recall: 0.8855 - val\_loss: 0.0288 - val\_acc: 0.9945 - val\_precision: 0.9952 - val\_recall: 0.9925  
Epoch 89/100  
137/137 [=====] - 1854s 14s/step - loss: 0.2618 - acc: 0.8994 - precision: 0.9118 - recall: 0.8863 - val\_loss: 0.0287 - val\_acc: 0.9953 - val\_precision: 0.9966 - val\_recall: 0.9946  
Epoch 90/100  
137/137 [=====] - 1859s 14s/step - loss: 0.2699 - acc: 0.9021 - precision: 0.9139 - recall: 0.8866 - val\_loss: 0.0294 - val\_acc: 0.9959 - val\_precision: 0.9959 - val\_recall: 0.9932  
Epoch 91/100  
137/137 [=====] - 1850s 14s/step - loss: 0.2615 - acc: 0.9036 - precision: 0.9138 - recall: 0.8895 - val\_loss: 0.0368 - val\_acc: 0.9925 - val\_precision: 0.9939 - val\_recall: 0.9912  
Epoch 92/100  
137/137 [=====] - 1589s 12s/step - loss: 0.2612 - acc: 0.9031 - precision: 0.9153 - recall: 0.8933 - val\_loss: 0.0340 - val\_acc: 0.9912 - val\_precision: 0.9912 - val\_recall: 0.9912  
Epoch 93/100  
137/137 [=====] - 1315s 10s/step - loss: 0.2473 - acc: 0.9112 - precision: 0.9215 - recall: 0.9001 - val\_loss: 0.0346 - val\_acc: 0.9919 - val\_precision: 0.9939 - val\_recall: 0.9912  
Epoch 94/100  
137/137 [=====] - 975s 7s/step - loss: 0.2502 - acc: 0.9112 - precision: 0.9210 - recall: 0.8985 - val\_loss: 0.0328 - val\_acc: 0.9878 - val\_precision: 0.9885 - val\_recall: 0.9878  
Epoch 95/100  
137/137 [=====] - 975s 7s/step - loss: 0.2416 - acc: 0.9153 - precision: 0.9243 - recall: 0.9041 - val\_loss: 0.0406 - val\_acc: 0.9870 - val\_precision: 0.9884 - val\_recall: 0.9857  
Epoch 96/100  
137/137 [=====] - 975s 7s/step - loss: 0.2531 - acc: 0.9116 - precision: 0.9228 - recall: 0.8995 - val\_loss: 0.0262 - val\_acc: 0.9939 - val\_precision: 0.9952 - val\_recall: 0.9932  
Epoch 97/100  
137/137 [=====] - 978s 7s/step - loss: 0.2317 - acc: 0.9174 - precision: 0.9288 - recall: 0.9078 - val\_loss: 0.0352 - val\_acc: 0.9878 - val\_precision: 0.9891 - val\_recall: 0.9865  
Epoch 98/100  
137/137 [=====] - 976s 7s/step - loss: 0.2612 - acc: 0.9049 - precision: 0.9184 - recall: 0.8935 - val\_loss: 0.0184 - val\_acc: 0.9980 - val\_precision: 0.9980 - val\_recall: 0.9966  
Epoch 99/100  
137/137 [=====] - 977s 7s/step - loss: 0.2463 - acc: 0.9094 - precision: 0.9220 - recall: 0.8990 - val\_loss: 0.0271 - val\_acc: 0.9939 - val\_precision: 0.9946 - val\_recall: 0.9926  
Epoch 100/100  
137/137 [=====] - 977s 7s/step - loss: 0.2331 - acc:

0.9155 - precision: 0.9275 - recall: 0.9050 - val\_loss: 0.0427 - val\_acc: 0.9  
 878 - val\_precision: 0.9884 - val\_recall: 0.9844

```
In [5]: test_steps_per_epoch = np.math.ceil(test_set.samples / test_set.batch_size)
        predictions = classifier.predict_generator(test_set, steps=test_steps_per_epoch)
        predicted_classes = np.argmax(predictions, axis=1)
```

```
In [6]: true_classes = test_set.classes
        class_labels = list(test_set.class_indices.keys())
```

```
In [7]: import sklearn.metrics as metrics
        report = metrics.classification_report(true_classes, predicted_classes, target_names=class_labels)
        print(report)
```

	precision	recall	f1-score	support
anger	0.23	0.23	0.23	350
boredom	0.19	0.19	0.19	223
disgust	0.08	0.08	0.08	130
fear	0.15	0.14	0.15	187
happiness	0.14	0.15	0.15	196
neutral	0.15	0.16	0.16	218
sadness	0.08	0.08	0.08	171
avg / total	0.16	0.16	0.16	1475

```

In [10]: import matplotlib.pyplot as plt
import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting normalize=True.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]*100
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap, aspect = 'auto')
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')


# Compute confusion matrix
cnf_matrix = metrics.confusion_matrix(true_classes, predicted_classes)
np.set_printoptions(precision=4)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_labels,
                      title='Confusion matrix, without normalization')
plt.savefig("non_normalized_confusion_matrix_cnn.png")
plt.show()
# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_labels, normalize=True,
                      title='Normalized confusion matrix')
plt.savefig("normalized_confusion_matrix_cnn.png")
plt.show()

```


Confusion matrix, without normalization

```
[[82 54 32 43 42 57 40]
 [58 42 22 24 32 29 16]
 [18 21 11 10 28 24 18]
 [50 25 16 27 18 28 23]
 [51 24 21 24 29 26 21]
 [49 20 16 28 31 34 40]
 [44 32 17 22 21 22 13]]
```



Normalized confusion matrix

```
[[23.4286 15.4286 9.1429 12.2857 12.0000 16.2857 11.4286]
 [26.009 18.8341 9.8655 10.7623 14.3498 13.0045 7.1749]
 [13.8462 16.1538 8.4615 7.6923 21.5385 18.4615 13.8462]
 [26.738 13.369 8.5561 14.4385 9.6257 14.9733 12.2995]
 [26.0204 12.2449 10.7143 12.2449 14.7959 13.2653 10.7143]
 [22.4771 9.1743 7.3394 12.844 14.2202 15.5963 18.3486]
 [25.731 18.7135 9.9415 12.8655 12.2807 12.8655 7.6023]]
```



```
In [11]: import matplotlib.pyplot as plt
plt.style.use("ggplot")
plt.figure()
N = 100
plt.plot(np.arange(0, N), results.history["loss"], label="train_loss")
plt.plot(np.arange(0, N), results.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, N), results.history["acc"], label="train_acc")
plt.plot(np.arange(0, N), results.history["val_acc"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend(loc="upper left")
plt.savefig("plot_cnn.png")
```

