

```
In [1]: # Convolutional Neural Network

# Installing Theano
# pip install --upgrade --no-deps git+git://github.com/Theano/Theano.git

# Installing Tensorflow
# Install Tensorflow from the website: https://www.tensorflow.org/versions/r0.12/get_started/os_setup.html

# Installing Keras
# pip install --upgrade keras

# Part 1 - Building the CNN

# Importing the Keras Libraries and packages
import numpy as np
import os
import keras_metrics
from keras.models import Sequential
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import TimeDistributed
from keras.layers import LSTM
from keras.layers import Reshape

import warnings
warnings.filterwarnings('ignore')

# Initialising the CNN
classifier = Sequential()

# Step 1 - Convolution
classifier.add(Convolution2D(64, (3, 3), padding = 'same', input_shape = (128, 128, 3), activation = 'relu'))

# Step 2 - Pooling
classifier.add(MaxPooling2D(pool_size = (2, 2)))

# Adding a second convolutional layer
classifier.add(Convolution2D(64, (3, 3), padding = 'same', activation = 'relu'))
classifier.add(MaxPooling2D(pool_size = (2, 2)))

# Adding a third convolutional layer
classifier.add(Convolution2D(64, (3, 3), padding = 'same', activation = 'relu'))
classifier.add(MaxPooling2D(pool_size = (2, 2)))

# Step 3 - Flattening
classifier.add(Flatten())
classifier.add(Dropout(rate = 0.5))

# Step 4 - Full connection
```

```

classifier.add(Reshape((4*4, 1024)))
classifier.add(LSTM(units = 50, return_sequences = True, dropout = 0.5))
classifier.add(LSTM(units = 20, return_sequences = False, dropout = 0.5))
classifier.add(Dense(output_dim = 7, activation = 'softmax'))

```

```

classifier.summary()

```

Z:\Anaconda3\lib\site-packages\h5py\\_\_init\_\_.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.

from .\_conv import register\_converters as \_register\_converters  
Using TensorFlow backend.

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 128, 128, 64)	1792
max_pooling2d_1 (MaxPooling2)	(None, 64, 64, 64)	0
conv2d_2 (Conv2D)	(None, 64, 64, 64)	36928
max_pooling2d_2 (MaxPooling2)	(None, 32, 32, 64)	0
conv2d_3 (Conv2D)	(None, 32, 32, 64)	36928
max_pooling2d_3 (MaxPooling2)	(None, 16, 16, 64)	0
flatten_1 (Flatten)	(None, 16384)	0
dropout_1 (Dropout)	(None, 16384)	0
reshape_1 (Reshape)	(None, 16, 1024)	0
lstm_1 (LSTM)	(None, 16, 50)	215000
lstm_2 (LSTM)	(None, 20)	5680
dense_1 (Dense)	(None, 7)	147
=====		
Total params: 296,475		
Trainable params: 296,475		
Non-trainable params: 0		
=====		

```

In [2]: # Compiling the CNN
classifier.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy', keras_metrics.precision(), keras_metrics.recall()])

```

```
In [3]: # Part 2 - Fitting the CNN to the images

from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(rescale = 1./255,
                                   shear_range = 0.2,
                                   zoom_range = 0.2,
                                   height_shift_range = 0.1,
                                   width_shift_range = 0.1,
                                   channel_shift_range = 10)

test_datagen = ImageDataGenerator(rescale = 1./255)

training_set = train_datagen.flow_from_directory('train/',
                                                target_size = (128, 128),
                                                batch_size = 32,
                                                class_mode = 'categorical')

test_set = test_datagen.flow_from_directory('test/',
                                            target_size = (128, 128),
                                            batch_size = 32,
                                            class_mode = 'categorical')
```

Found 3960 images belonging to 7 classes.  
Found 1320 images belonging to 7 classes.

```
In [4]: results = classifier.fit_generator(training_set,
                                         samples_per_epoch = 3960,
                                         nb_epoch = 100,
                                         validation_data = test_set,
                                         nb_val_samples = 1320)
```

Epoch 1/100  
123/123 [=====] - 1666s 14s/step - loss: 1.8975 - acc: 0.2580 - precision: 0.0000e+00 - recall: 0.0000e+00 - val\_loss: 1.7336 - val\_acc: 0.3042 - val\_precision: 0.0000e+00 - val\_recall: 0.0000e+00

Epoch 2/100  
123/123 [=====] - 1670s 14s/step - loss: 1.7140 - acc: 0.3191 - precision: 0.4906 - recall: 0.0716 - val\_loss: 1.5569 - val\_acc: 0.3510 - val\_precision: 0.6223 - val\_recall: 0.1023

Epoch 3/100  
123/123 [=====] - 1672s 14s/step - loss: 1.6481 - acc: 0.3218 - precision: 0.5623 - recall: 0.1017 - val\_loss: 1.5345 - val\_acc: 0.3741 - val\_precision: 0.5933 - val\_recall: 0.0499

Epoch 4/100  
123/123 [=====] - 1669s 14s/step - loss: 1.6069 - acc: 0.3444 - precision: 0.5926 - recall: 0.1275 - val\_loss: 1.4606 - val\_acc: 0.3865 - val\_precision: 0.6224 - val\_recall: 0.1804

Epoch 5/100  
123/123 [=====] - 1677s 14s/step - loss: 1.5628 - acc: 0.3683 - precision: 0.6108 - recall: 0.1276 - val\_loss: 1.4638 - val\_acc: 0.3999 - val\_precision: 0.6296 - val\_recall: 0.1475

Epoch 6/100  
123/123 [=====] - 1670s 14s/step - loss: 1.5411 - acc: 0.3686 - precision: 0.6092 - recall: 0.1344 - val\_loss: 1.3739 - val\_acc: 0.4275 - val\_precision: 0.6538 - val\_recall: 0.1932

Epoch 7/100  
123/123 [=====] - 1674s 14s/step - loss: 1.5198 - acc: 0.3836 - precision: 0.6268 - recall: 0.1358 - val\_loss: 1.4771 - val\_acc: 0.3910 - val\_precision: 0.6984 - val\_recall: 0.1402

Epoch 8/100  
123/123 [=====] - 1678s 14s/step - loss: 1.4745 - acc: 0.4098 - precision: 0.6496 - recall: 0.1542 - val\_loss: 1.4014 - val\_acc: 0.4391 - val\_precision: 0.6858 - val\_recall: 0.1999

Epoch 9/100  
123/123 [=====] - 1682s 14s/step - loss: 1.4144 - acc: 0.4350 - precision: 0.6569 - recall: 0.1875 - val\_loss: 1.2835 - val\_acc: 0.4710 - val\_precision: 0.7474 - val\_recall: 0.1798

Epoch 10/100  
123/123 [=====] - 1669s 14s/step - loss: 1.3772 - acc: 0.4572 - precision: 0.6718 - recall: 0.2060 - val\_loss: 1.1897 - val\_acc: 0.5283 - val\_precision: 0.6874 - val\_recall: 0.2941

Epoch 11/100  
123/123 [=====] - 1673s 14s/step - loss: 1.3651 - acc: 0.4608 - precision: 0.6655 - recall: 0.2047 - val\_loss: 1.1196 - val\_acc: 0.5449 - val\_precision: 0.7172 - val\_recall: 0.3262

Epoch 12/100  
123/123 [=====] - 1715s 14s/step - loss: 1.3230 - acc: 0.4813 - precision: 0.6507 - recall: 0.2417 - val\_loss: 1.0873 - val\_acc: 0.5634 - val\_precision: 0.7644 - val\_recall: 0.3286

Epoch 13/100  
123/123 [=====] - 1681s 14s/step - loss: 1.2813 - acc: 0.5005 - precision: 0.6888 - recall: 0.2632 - val\_loss: 1.0763 - val\_acc: 0.5716 - val\_precision: 0.7008 - val\_recall: 0.3480

Epoch 14/100  
123/123 [=====] - 1676s 14s/step - loss: 1.2447 - acc: 0.5142 - precision: 0.6901 - recall: 0.2890 - val\_loss: 1.0099 - val\_acc: 0.5889 - val\_precision: 0.7311 - val\_recall: 0.3971

Epoch 15/100

```
123/123 [=====] - 1676s 14s/step - loss: 1.2238 - ac
c: 0.5163 - precision: 0.6941 - recall: 0.2993 - val_loss: 1.0212 - val_acc:
0.5864 - val_precision: 0.7078 - val_recall: 0.4060
Epoch 16/100
123/123 [=====] - 1675s 14s/step - loss: 1.2081 - ac
c: 0.5223 - precision: 0.6909 - recall: 0.3073 - val_loss: 1.0105 - val_acc:
0.5845 - val_precision: 0.7217 - val_recall: 0.4120
Epoch 17/100
123/123 [=====] - 1679s 14s/step - loss: 1.1917 - ac
c: 0.5326 - precision: 0.6888 - recall: 0.3319 - val_loss: 0.9666 - val_acc:
0.6100 - val_precision: 0.7074 - val_recall: 0.4440
Epoch 18/100
123/123 [=====] - 1672s 14s/step - loss: 1.1808 - ac
c: 0.5309 - precision: 0.6942 - recall: 0.3336 - val_loss: 0.9067 - val_acc:
0.6534 - val_precision: 0.7742 - val_recall: 0.4338
Epoch 19/100
123/123 [=====] - 1683s 14s/step - loss: 1.1426 - ac
c: 0.5569 - precision: 0.7059 - recall: 0.3534 - val_loss: 0.9754 - val_acc:
0.6182 - val_precision: 0.6947 - val_recall: 0.4509
Epoch 20/100
123/123 [=====] - 1676s 14s/step - loss: 1.0940 - ac
c: 0.5688 - precision: 0.7127 - recall: 0.3747 - val_loss: 0.8945 - val_acc:
0.6696 - val_precision: 0.7591 - val_recall: 0.4915
Epoch 21/100
123/123 [=====] - 1678s 14s/step - loss: 1.1065 - ac
c: 0.5650 - precision: 0.7046 - recall: 0.3808 - val_loss: 0.8344 - val_acc:
0.6947 - val_precision: 0.7676 - val_recall: 0.5364
Epoch 22/100
123/123 [=====] - 1677s 14s/step - loss: 1.0971 - ac
c: 0.5777 - precision: 0.6963 - recall: 0.3877 - val_loss: 0.9292 - val_acc:
0.6362 - val_precision: 0.7054 - val_recall: 0.5159
Epoch 23/100
123/123 [=====] - 1677s 14s/step - loss: 1.0447 - ac
c: 0.5815 - precision: 0.7123 - recall: 0.4201 - val_loss: 0.8250 - val_acc:
0.6706 - val_precision: 0.7393 - val_recall: 0.5668
Epoch 24/100
123/123 [=====] - 1675s 14s/step - loss: 1.0335 - ac
c: 0.6069 - precision: 0.7146 - recall: 0.4432 - val_loss: 0.7901 - val_acc:
0.6885 - val_precision: 0.7855 - val_recall: 0.5778
Epoch 25/100
123/123 [=====] - 1673s 14s/step - loss: 1.0108 - ac
c: 0.6026 - precision: 0.7149 - recall: 0.4568 - val_loss: 0.8481 - val_acc:
0.6545 - val_precision: 0.7267 - val_recall: 0.5629
Epoch 26/100
123/123 [=====] - 1675s 14s/step - loss: 1.0105 - ac
c: 0.6088 - precision: 0.7137 - recall: 0.4498 - val_loss: 0.7485 - val_acc:
0.7131 - val_precision: 0.7694 - val_recall: 0.6303
Epoch 27/100
123/123 [=====] - 1673s 14s/step - loss: 0.9844 - ac
c: 0.6237 - precision: 0.7269 - recall: 0.4747 - val_loss: 0.7266 - val_acc:
0.7263 - val_precision: 0.7937 - val_recall: 0.6375
Epoch 28/100
123/123 [=====] - 1695s 14s/step - loss: 0.9823 - ac
c: 0.6249 - precision: 0.7201 - recall: 0.4853 - val_loss: 0.6933 - val_acc:
0.7385 - val_precision: 0.7955 - val_recall: 0.6506
Epoch 29/100
123/123 [=====] - 1690s 14s/step - loss: 0.9346 - ac
```

c: 0.6456 - precision: 0.7389 - recall: 0.5159 - val\_loss: 0.6852 - val\_acc: 0.7446 - val\_precision: 0.7992 - val\_recall: 0.6762  
Epoch 30/100  
123/123 [=====] - 1676s 14s/step - loss: 0.9200 - acc: 0.6413 - precision: 0.7431 - recall: 0.5288 - val\_loss: 0.6326 - val\_acc: 0.7771 - val\_precision: 0.8355 - val\_recall: 0.7088  
Epoch 31/100  
123/123 [=====] - 1675s 14s/step - loss: 0.8975 - acc: 0.6625 - precision: 0.7534 - recall: 0.5407 - val\_loss: 0.6162 - val\_acc: 0.7855 - val\_precision: 0.8283 - val\_recall: 0.7128  
Epoch 32/100  
123/123 [=====] - 1674s 14s/step - loss: 0.8680 - acc: 0.6761 - precision: 0.7501 - recall: 0.5617 - val\_loss: 0.5964 - val\_acc: 0.7870 - val\_precision: 0.8270 - val\_recall: 0.7386  
Epoch 33/100  
123/123 [=====] - 1680s 14s/step - loss: 0.8885 - acc: 0.6585 - precision: 0.7393 - recall: 0.5567 - val\_loss: 0.6459 - val\_acc: 0.7696 - val\_precision: 0.8205 - val\_recall: 0.6968  
Epoch 34/100  
123/123 [=====] - 1672s 14s/step - loss: 0.8477 - acc: 0.6850 - precision: 0.7515 - recall: 0.5741 - val\_loss: 0.5449 - val\_acc: 0.8025 - val\_precision: 0.8507 - val\_recall: 0.7502  
Epoch 35/100  
123/123 [=====] - 1685s 14s/step - loss: 0.8479 - acc: 0.6773 - precision: 0.7548 - recall: 0.5823 - val\_loss: 0.5331 - val\_acc: 0.8082 - val\_precision: 0.8576 - val\_recall: 0.7627  
Epoch 36/100  
123/123 [=====] - 1678s 14s/step - loss: 0.8072 - acc: 0.6988 - precision: 0.7735 - recall: 0.6129 - val\_loss: 0.5356 - val\_acc: 0.8024 - val\_precision: 0.8367 - val\_recall: 0.7674  
Epoch 37/100  
123/123 [=====] - 1677s 14s/step - loss: 0.7833 - acc: 0.7089 - precision: 0.7728 - recall: 0.6231 - val\_loss: 0.5720 - val\_acc: 0.7850 - val\_precision: 0.8145 - val\_recall: 0.7448  
Epoch 38/100  
123/123 [=====] - 1677s 14s/step - loss: 0.7954 - acc: 0.6961 - precision: 0.7659 - recall: 0.6143 - val\_loss: 0.4872 - val\_acc: 0.8220 - val\_precision: 0.8513 - val\_recall: 0.7756  
Epoch 39/100  
123/123 [=====] - 1670s 14s/step - loss: 0.7701 - acc: 0.7051 - precision: 0.7691 - recall: 0.6306 - val\_loss: 0.5657 - val\_acc: 0.7870 - val\_precision: 0.8198 - val\_recall: 0.7537  
Epoch 40/100  
123/123 [=====] - 1675s 14s/step - loss: 0.7414 - acc: 0.7225 - precision: 0.7898 - recall: 0.6542 - val\_loss: 0.5262 - val\_acc: 0.8136 - val\_precision: 0.8371 - val\_recall: 0.7818  
Epoch 41/100  
123/123 [=====] - 1679s 14s/step - loss: 0.7378 - acc: 0.7261 - precision: 0.7846 - recall: 0.6552 - val\_loss: 0.5114 - val\_acc: 0.8157 - val\_precision: 0.8440 - val\_recall: 0.7838  
Epoch 42/100  
123/123 [=====] - 1673s 14s/step - loss: 0.7470 - acc: 0.7264 - precision: 0.7838 - recall: 0.6498 - val\_loss: 0.4207 - val\_acc: 0.8490 - val\_precision: 0.8806 - val\_recall: 0.8262  
Epoch 43/100  
123/123 [=====] - 1675s 14s/step - loss: 0.7001 - acc: 0.7381 - precision: 0.7900 - recall: 0.6803 - val\_loss: 0.6045 - val\_acc:

0.7741 - val\_precision: 0.8177 - val\_recall: 0.7452  
Epoch 44/100  
123/123 [=====] - 1679s 14s/step - loss: 0.7222 - acc: 0.7371 - precision: 0.7917 - recall: 0.6709 - val\_loss: 0.4157 - val\_acc: 0.8470 - val\_precision: 0.8794 - val\_recall: 0.8227  
Epoch 45/100  
123/123 [=====] - 1681s 14s/step - loss: 0.6859 - acc: 0.7491 - precision: 0.8002 - recall: 0.6878 - val\_loss: 0.3928 - val\_acc: 0.8554 - val\_precision: 0.8717 - val\_recall: 0.8274  
Epoch 46/100  
123/123 [=====] - 1733s 14s/step - loss: 0.7016 - acc: 0.7416 - precision: 0.7947 - recall: 0.6855 - val\_loss: 0.4112 - val\_acc: 0.8561 - val\_precision: 0.8742 - val\_recall: 0.8213  
Epoch 47/100  
123/123 [=====] - 1718s 14s/step - loss: 0.6765 - acc: 0.7455 - precision: 0.7986 - recall: 0.6915 - val\_loss: 0.3821 - val\_acc: 0.8672 - val\_precision: 0.8870 - val\_recall: 0.8520  
Epoch 48/100  
123/123 [=====] - 1682s 14s/step - loss: 0.6390 - acc: 0.7708 - precision: 0.8211 - recall: 0.7231 - val\_loss: 0.3519 - val\_acc: 0.8824 - val\_precision: 0.8938 - val\_recall: 0.8627  
Epoch 49/100  
123/123 [=====] - 1685s 14s/step - loss: 0.6364 - acc: 0.7675 - precision: 0.8091 - recall: 0.7149 - val\_loss: 0.3583 - val\_acc: 0.8828 - val\_precision: 0.9007 - val\_recall: 0.8637  
Epoch 50/100  
123/123 [=====] - 1679s 14s/step - loss: 0.6540 - acc: 0.7588 - precision: 0.8085 - recall: 0.7108 - val\_loss: 0.4984 - val\_acc: 0.8260 - val\_precision: 0.8435 - val\_recall: 0.7996  
Epoch 51/100  
123/123 [=====] - 1686s 14s/step - loss: 0.6113 - acc: 0.7767 - precision: 0.8206 - recall: 0.7264 - val\_loss: 0.3744 - val\_acc: 0.8615 - val\_precision: 0.8827 - val\_recall: 0.8471  
Epoch 52/100  
123/123 [=====] - 1685s 14s/step - loss: 0.5902 - acc: 0.7917 - precision: 0.8335 - recall: 0.7447 - val\_loss: 0.3151 - val\_acc: 0.8997 - val\_precision: 0.9128 - val\_recall: 0.8837  
Epoch 53/100  
123/123 [=====] - 1691s 14s/step - loss: 0.6148 - acc: 0.7797 - precision: 0.8196 - recall: 0.7284 - val\_loss: 0.3271 - val\_acc: 0.8895 - val\_precision: 0.9062 - val\_recall: 0.8781  
Epoch 54/100  
123/123 [=====] - 1689s 14s/step - loss: 0.5530 - acc: 0.7982 - precision: 0.8376 - recall: 0.7595 - val\_loss: 0.3351 - val\_acc: 0.8932 - val\_precision: 0.9020 - val\_recall: 0.8780  
Epoch 55/100  
123/123 [=====] - 1686s 14s/step - loss: 0.5757 - acc: 0.7916 - precision: 0.8287 - recall: 0.7511 - val\_loss: 0.3037 - val\_acc: 0.8986 - val\_precision: 0.9129 - val\_recall: 0.8812  
Epoch 56/100  
123/123 [=====] - 1682s 14s/step - loss: 0.5849 - acc: 0.7843 - precision: 0.8260 - recall: 0.7416 - val\_loss: 0.2821 - val\_acc: 0.9060 - val\_precision: 0.9187 - val\_recall: 0.8909  
Epoch 57/100  
123/123 [=====] - 1689s 14s/step - loss: 0.5486 - acc: 0.8014 - precision: 0.8382 - recall: 0.7653 - val\_loss: 0.2798 - val\_acc: 0.9051 - val\_precision: 0.9150 - val\_recall: 0.8908



Epoch 58/100  
123/123 [=====] - 1683s 14s/step - loss: 0.5425 - acc: 0.8095 - precision: 0.8425 - recall: 0.7717 - val\_loss: 0.2770 - val\_acc: 0.9061 - val\_precision: 0.9168 - val\_recall: 0.8939

Epoch 59/100  
123/123 [=====] - 1699s 14s/step - loss: 0.5336 - acc: 0.8057 - precision: 0.8482 - recall: 0.7727 - val\_loss: 0.2178 - val\_acc: 0.9318 - val\_precision: 0.9449 - val\_recall: 0.9211

Epoch 60/100  
123/123 [=====] - 1684s 14s/step - loss: 0.5189 - acc: 0.8150 - precision: 0.8462 - recall: 0.7782 - val\_loss: 0.2568 - val\_acc: 0.9137 - val\_precision: 0.9191 - val\_recall: 0.9024

Epoch 61/100  
123/123 [=====] - 1677s 14s/step - loss: 0.5114 - acc: 0.8209 - precision: 0.8528 - recall: 0.7834 - val\_loss: 0.2567 - val\_acc: 0.9141 - val\_precision: 0.9265 - val\_recall: 0.8997

Epoch 62/100  
123/123 [=====] - 1681s 14s/step - loss: 0.5034 - acc: 0.8196 - precision: 0.8562 - recall: 0.7908 - val\_loss: 0.2628 - val\_acc: 0.9070 - val\_precision: 0.9222 - val\_recall: 0.8964

Epoch 63/100  
123/123 [=====] - 1691s 14s/step - loss: 0.5045 - acc: 0.8226 - precision: 0.8532 - recall: 0.7892 - val\_loss: 0.3107 - val\_acc: 0.8970 - val\_precision: 0.9146 - val\_recall: 0.8834

Epoch 64/100  
123/123 [=====] - 1679s 14s/step - loss: 0.5119 - acc: 0.8116 - precision: 0.8467 - recall: 0.7798 - val\_loss: 0.1882 - val\_acc: 0.9394 - val\_precision: 0.9520 - val\_recall: 0.9311

Epoch 65/100  
123/123 [=====] - 1678s 14s/step - loss: 0.4866 - acc: 0.8298 - precision: 0.8593 - recall: 0.7975 - val\_loss: 0.2079 - val\_acc: 0.9425 - val\_precision: 0.9479 - val\_recall: 0.9365

Epoch 66/100  
123/123 [=====] - 1691s 14s/step - loss: 0.4794 - acc: 0.8238 - precision: 0.8524 - recall: 0.7939 - val\_loss: 0.2677 - val\_acc: 0.9061 - val\_precision: 0.9144 - val\_recall: 0.8971

Epoch 67/100  
123/123 [=====] - 1691s 14s/step - loss: 0.4515 - acc: 0.8357 - precision: 0.8628 - recall: 0.8142 - val\_loss: 0.2130 - val\_acc: 0.9342 - val\_precision: 0.9397 - val\_recall: 0.9311

Epoch 68/100  
123/123 [=====] - 1689s 14s/step - loss: 0.4654 - acc: 0.8387 - precision: 0.8673 - recall: 0.8100 - val\_loss: 0.1977 - val\_acc: 0.9439 - val\_precision: 0.9516 - val\_recall: 0.9385

Epoch 69/100  
123/123 [=====] - 1684s 14s/step - loss: 0.4689 - acc: 0.8368 - precision: 0.8662 - recall: 0.8058 - val\_loss: 0.1936 - val\_acc: 0.9394 - val\_precision: 0.9492 - val\_recall: 0.9326

Epoch 70/100  
123/123 [=====] - 1679s 14s/step - loss: 0.4635 - acc: 0.8328 - precision: 0.8621 - recall: 0.8041 - val\_loss: 0.1619 - val\_acc: 0.9516 - val\_precision: 0.9607 - val\_recall: 0.9447

Epoch 71/100  
123/123 [=====] - 1691s 14s/step - loss: 0.4709 - acc: 0.8357 - precision: 0.8651 - recall: 0.8074 - val\_loss: 0.1757 - val\_acc: 0.9454 - val\_precision: 0.9531 - val\_recall: 0.9393

Epoch 72/100

```
123/123 [=====] - 1690s 14s/step - loss: 0.4305 - ac
c: 0.8486 - precision: 0.8759 - recall: 0.8229 - val_loss: 0.1584 - val_acc:
0.9562 - val_precision: 0.9652 - val_recall: 0.9463
Epoch 73/100
123/123 [=====] - 1688s 14s/step - loss: 0.4378 - ac
c: 0.8466 - precision: 0.8720 - recall: 0.8225 - val_loss: 0.1627 - val_acc:
0.9575 - val_precision: 0.9618 - val_recall: 0.9552
Epoch 74/100
123/123 [=====] - 1681s 14s/step - loss: 0.4352 - ac
c: 0.8442 - precision: 0.8692 - recall: 0.8173 - val_loss: 0.2960 - val_acc:
0.8969 - val_precision: 0.9078 - val_recall: 0.8878
Epoch 75/100
123/123 [=====] - 1689s 14s/step - loss: 0.4308 - ac
c: 0.8479 - precision: 0.8693 - recall: 0.8258 - val_loss: 0.1320 - val_acc:
0.9652 - val_precision: 0.9709 - val_recall: 0.9621
Epoch 76/100
123/123 [=====] - 1692s 14s/step - loss: 0.4168 - ac
c: 0.8540 - precision: 0.8809 - recall: 0.8314 - val_loss: 0.2131 - val_acc:
0.9368 - val_precision: 0.9438 - val_recall: 0.9331
Epoch 77/100
123/123 [=====] - 1685s 14s/step - loss: 0.4271 - ac
c: 0.8497 - precision: 0.8756 - recall: 0.8263 - val_loss: 0.1846 - val_acc:
0.9447 - val_precision: 0.9480 - val_recall: 0.9372
Epoch 78/100
123/123 [=====] - 1686s 14s/step - loss: 0.3865 - ac
c: 0.8695 - precision: 0.8881 - recall: 0.8492 - val_loss: 0.1598 - val_acc:
0.9476 - val_precision: 0.9548 - val_recall: 0.9468
Epoch 79/100
123/123 [=====] - 1687s 14s/step - loss: 0.3946 - ac
c: 0.8685 - precision: 0.8852 - recall: 0.8460 - val_loss: 0.1359 - val_acc:
0.9622 - val_precision: 0.9665 - val_recall: 0.9584
Epoch 80/100
123/123 [=====] - 1685s 14s/step - loss: 0.3827 - ac
c: 0.8684 - precision: 0.8901 - recall: 0.8514 - val_loss: 0.1578 - val_acc:
0.9545 - val_precision: 0.9588 - val_recall: 0.9522
Epoch 81/100
123/123 [=====] - 1699s 14s/step - loss: 0.3935 - ac
c: 0.8653 - precision: 0.8896 - recall: 0.8448 - val_loss: 0.1533 - val_acc:
0.9576 - val_precision: 0.9670 - val_recall: 0.9546
Epoch 82/100
123/123 [=====] - 1701s 14s/step - loss: 0.3568 - ac
c: 0.8775 - precision: 0.8986 - recall: 0.8618 - val_loss: 0.1468 - val_acc:
0.9560 - val_precision: 0.9616 - val_recall: 0.9485
Epoch 83/100
123/123 [=====] - 1698s 14s/step - loss: 0.3716 - ac
c: 0.8717 - precision: 0.8923 - recall: 0.8528 - val_loss: 0.1150 - val_acc:
0.9659 - val_precision: 0.9694 - val_recall: 0.9613
Epoch 84/100
123/123 [=====] - 1695s 14s/step - loss: 0.3844 - ac
c: 0.8712 - precision: 0.8930 - recall: 0.8542 - val_loss: 0.1057 - val_acc:
0.9719 - val_precision: 0.9741 - val_recall: 0.9704
Epoch 85/100
123/123 [=====] - 1695s 14s/step - loss: 0.3383 - ac
c: 0.8836 - precision: 0.9015 - recall: 0.8653 - val_loss: 0.1041 - val_acc:
0.9705 - val_precision: 0.9727 - val_recall: 0.9705
Epoch 86/100
123/123 [=====] - 1687s 14s/step - loss: 0.3334 - ac
```

c: 0.8887 - precision: 0.9050 - recall: 0.8742 - val\_loss: 0.0998 - val\_acc: 0.9758 - val\_precision: 0.9772 - val\_recall: 0.9735  
Epoch 87/100  
123/123 [=====] - 1701s 14s/step - loss: 0.3576 - acc: 0.8780 - precision: 0.8913 - recall: 0.8645 - val\_loss: 0.0911 - val\_acc: 0.9773 - val\_precision: 0.9802 - val\_recall: 0.9735  
Epoch 88/100  
123/123 [=====] - 1710s 14s/step - loss: 0.3693 - acc: 0.8743 - precision: 0.8920 - recall: 0.8583 - val\_loss: 0.1013 - val\_acc: 0.9727 - val\_precision: 0.9771 - val\_recall: 0.9705  
Epoch 89/100  
123/123 [=====] - 1715s 14s/step - loss: 0.3542 - acc: 0.8852 - precision: 0.9012 - recall: 0.8689 - val\_loss: 0.1415 - val\_acc: 0.9507 - val\_precision: 0.9563 - val\_recall: 0.9455  
Epoch 90/100  
123/123 [=====] - 1701s 14s/step - loss: 0.3326 - acc: 0.8893 - precision: 0.9044 - recall: 0.8725 - val\_loss: 0.1034 - val\_acc: 0.9713 - val\_precision: 0.9734 - val\_recall: 0.9690  
Epoch 91/100  
123/123 [=====] - 1699s 14s/step - loss: 0.3490 - acc: 0.8789 - precision: 0.8953 - recall: 0.8620 - val\_loss: 0.0990 - val\_acc: 0.9713 - val\_precision: 0.9727 - val\_recall: 0.9675  
Epoch 92/100  
123/123 [=====] - 1694s 14s/step - loss: 0.3505 - acc: 0.8817 - precision: 0.9007 - recall: 0.8667 - val\_loss: 0.1108 - val\_acc: 0.9658 - val\_precision: 0.9701 - val\_recall: 0.9636  
Epoch 93/100  
123/123 [=====] - 1694s 14s/step - loss: 0.3436 - acc: 0.8854 - precision: 0.9032 - recall: 0.8684 - val\_loss: 0.1275 - val\_acc: 0.9628 - val\_precision: 0.9657 - val\_recall: 0.9598  
Epoch 94/100  
123/123 [=====] - 1691s 14s/step - loss: 0.3087 - acc: 0.8984 - precision: 0.9138 - recall: 0.8814 - val\_loss: 0.1100 - val\_acc: 0.9696 - val\_precision: 0.9725 - val\_recall: 0.9666  
Epoch 95/100  
123/123 [=====] - 1702s 14s/step - loss: 0.3210 - acc: 0.8877 - precision: 0.9053 - recall: 0.8703 - val\_loss: 0.0928 - val\_acc: 0.9742 - val\_precision: 0.9757 - val\_recall: 0.9727  
Epoch 96/100  
123/123 [=====] - 1697s 14s/step - loss: 0.3296 - acc: 0.8890 - precision: 0.9023 - recall: 0.8725 - val\_loss: 0.0899 - val\_acc: 0.9750 - val\_precision: 0.9772 - val\_recall: 0.9735  
Epoch 97/100  
123/123 [=====] - 1699s 14s/step - loss: 0.3179 - acc: 0.8930 - precision: 0.9106 - recall: 0.8787 - val\_loss: 0.1057 - val\_acc: 0.9696 - val\_precision: 0.9726 - val\_recall: 0.9674  
Epoch 98/100  
123/123 [=====] - 1682s 14s/step - loss: 0.3051 - acc: 0.8956 - precision: 0.9089 - recall: 0.8801 - val\_loss: 0.1046 - val\_acc: 0.9720 - val\_precision: 0.9741 - val\_recall: 0.9705  
Epoch 99/100  
123/123 [=====] - 1477s 12s/step - loss: 0.3008 - acc: 0.8941 - precision: 0.9073 - recall: 0.8852 - val\_loss: 0.1150 - val\_acc: 0.9628 - val\_precision: 0.9664 - val\_recall: 0.9598  
Epoch 100/100  
123/123 [=====] - 1281s 10s/step - loss: 0.2955 - acc:

c: 0.8996 - precision: 0.9140 - recall: 0.8877 - val\_loss: 0.0721 - val\_acc: 0.9810 - val\_precision: 0.9832 - val\_recall: 0.9788

```
In [5]: test_steps_per_epoch = np.math.ceil(test_set.samples / test_set.batch_size)
        predictions = classifier.predict_generator(test_set, steps=test_steps_per_epoch)
        predicted_classes = np.argmax(predictions, axis=1)
```

```
In [6]: true_classes = test_set.classes
        class_labels = list(test_set.class_indices.keys())
```

```
In [7]: import sklearn.metrics as metrics
        report = metrics.classification_report(true_classes, predicted_classes, target_names=class_labels)
        print(report)
```

	precision	recall	f1-score	support
angry	0.13	0.13	0.13	165
disgust	0.09	0.08	0.09	165
fearful	0.13	0.13	0.13	165
happy	0.12	0.12	0.12	165
neutral	0.23	0.24	0.24	330
sad	0.16	0.16	0.16	165
surprised	0.12	0.12	0.12	165
avg / total	0.15	0.15	0.15	1320

```

In [10]: import matplotlib.pyplot as plt
import itertools
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting normalize=True.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]*100
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap, aspect = 'auto')
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')


# Compute confusion matrix
cnf_matrix = metrics.confusion_matrix(true_classes, predicted_classes)
np.set_printoptions(precision=4)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_labels,
                      title='Confusion matrix, without normalization')
plt.savefig("non_normalized_confusion_matrix_cnn_lstm.png")
plt.show()
# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=class_labels, normalize=True,
                      title='Normalized confusion matrix')
plt.savefig("normalized_confusion_matrix_cnn_lstm.png")
plt.show()

```


Confusion matrix, without normalization

```
[[22 21 18 21 43 21 19]
 [11 14 23 27 52 17 21]
 [23 26 22 17 41 18 18]
 [28 18 19 20 41 16 23]
 [43 38 40 46 79 45 39]
 [23 16 17 17 43 26 23]
 [16 28 24 22 38 17 20]]
```



Normalized confusion matrix

```
[[13.3333 12.7273 10.9091 12.7273 26.0606 12.7273 11.5152]
 [ 6.6667  8.4848 13.9394 16.3636 31.5152 10.303 12.7273]
 [13.9394 15.7576 13.3333 10.303 24.8485 10.9091 10.9091]
 [16.9697 10.9091 11.5152 12.1212 24.8485  9.697 13.9394]
 [13.0303 11.5152 12.1212 13.9394 23.9394 13.6364 11.8182]
 [13.9394  9.697 10.303 10.303 26.0606 15.7576 13.9394]
 [ 9.697 16.9697 14.5455 13.3333 23.0303 10.303 12.1212]]
```



```
In [11]: import matplotlib.pyplot as plt
plt.style.use("ggplot")
plt.figure()
N = 100
plt.plot(np.arange(0, N), results.history["loss"], label="train_loss")
plt.plot(np.arange(0, N), results.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, N), results.history["acc"], label="train_acc")
plt.plot(np.arange(0, N), results.history["val_acc"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend(loc="upper left")
plt.savefig("plot_cnn_lstm.png")
```

